

Computing for Analytical Work

Session 3 — What broke?

Block 5 — Reading errors and debugging systematically

Knowledge check 3 — 10 minutes, closed book, 3 questions

- Covers **Session 2 and Homework 2**: interpreters, environments, `uv`, notebook state, Git checkpoints
- Short answers, in your own words
- Diagnostic, not tricky — if you did the work, this is easy
- Closed book means closed: no laptop, no phone, no notes
- **This block's lecture is 35 minutes**, not 45 — the check comes out of lecture, never out of lab
- Hand it in when the timer ends; the lecture starts immediately after

This block in one sentence

Error messages are evidence. Debugging is a method, not a vibe.

A traceback is not a punishment. It is the machine telling you, precisely, what it could not do and where it gave up.

Agenda

1. What a traceback is, and the two anchors
2. Symptom versus cause
3. Six errors you will meet this year
4. Inspect before you guess: print, minimal example, asking well
5. Git thread — `git log` : what happened before?
6. AI thread — a prompt made of evidence
7. Lab 5 — The traceback is the map

What a traceback is

- Python hit something it could not do, and stopped
- On its way out, it wrote down the **chain of calls** that led there
- It is printed **top to bottom**: the outermost call first, the failure last
- That is why the useful part is at the **bottom**
- Nothing about it is random; every line points at a file and a line number

First: find it in the scroll

- The traceback begins at the literal line `Traceback (most recent call last):` — that is what you search for
- Scroll back: trackpad, scrollbar, or **Cmd-F / Ctrl-F** in VS Code's terminal
- Several runs' output stacked up? You cannot tell which error is *this* run's
- The habit that removes the problem: **clear** before every rerun — then everything on screen is from this run

The two anchors

1. **Last line** → the **error type** → *what kind of problem*
2. **Nearest line above that is in YOUR file** → the **failing line** → *where*

Everything else is the call stack. Scroll up later, if the two anchors are not enough.

A real one

```
Traceback (most recent call last):
  File "/Users/anna/lab05/scripts/report.py", line 14, in <module>
    total = monthly_total(df)
  File "/Users/anna/lab05/scripts/report.py", line 8, in monthly_total
    return df["revenue"].sum()
  File ".../site-packages/pandas/core/frame.py", line 4102, in __getitem__
    indexer = self.columns.get_loc(key)
  File ".../site-packages/pandas/core/indexes/base.py", line 3812, in get_loc
    raise KeyError(key) from err
KeyError: 'revenue'
```

The same traceback, annotated

```

File ".../scripts/report.py", line 8, in monthly_total
    return df["revenue"].sum()
File ".../site-packages/pandas/core/frame.py" ...
File ".../site-packages/pandas/core/indexes/base.py" ...
KeyError: 'revenue'

```

<- ANCHOR 2: where
 <- the failing line
 <- call stack (pandas)
 <- call stack (pandas)
 <- ANCHOR 1: what

- **What:** `KeyError` — a name was looked up and did not exist
- **Where:** `report.py` line 8, `df["revenue"]`
- **Next:** `print(df.columns)` — what names *does* the dataframe have?

When the library is in the way

- Frames in `site-packages/` belong to pandas, numpy, matplotlib — **not to you**
- The bug is almost never inside the library
- Walk **up** from the bottom until the path is a file in **your** project
- That frame is anchor 2, even if it is four lines above the error
- Notebook tracebacks look different but obey the same rule: find the cell, find the line

Symptom versus cause

- **Symptom:** what the traceback says — `KeyError: 'revenue'`
- **Cause:** why the world is that way — the CSV header says `Revenue`, capital R
- The traceback gives you the symptom for free; the cause you have to go find
- Fixing the symptom: `df["Revenue"]` — works until the next file
- Fixing the cause: normalise column names right after loading
- The diagnosis note asks for both, on purpose

FileNotFoundError and **ModuleNotFoundError**

FileNotFoundError: [Errno 2] No such file or directory: 'data/sales.csv'

- Usually means: the path is relative to the **wrong working directory**, or the file is in `data/raw/`
- Inspect first: `pwd` , `ls data/` , `Path.cwd()` in the notebook

ModuleNotFoundError: No module named 'pandas'

- Usually means: you are running a **different Python** than the one you installed into
- Inspect first: `uv run python -c "import sys; print(sys.executable)"`

NameError and **KeyError**

NameError: name 'df' is not defined

- Usually means: the line that creates `df` never ran — typo, skipped cell, wrong order
- Inspect first: search the file for `df =`; in a notebook, Restart and Run All

KeyError: 'revenue'

- Usually means: a column, dictionary key, or config entry is **spelled differently** than you think
- Inspect first: `print(df.columns.tolist())` or `print(d.keys())`

ValueError and TypeError

ValueError: could not convert string to float: '€1,200'

- Usually means: the **type is right, the content is wrong** — a number with a currency sign, a date with the wrong format
- Inspect first: look at the actual values — `df["amount"].head()`, `df["amount"].unique()`
`[:10]`

TypeError: unsupported operand type(s) for +: 'int' and 'str'

- Usually means: you passed the **wrong kind of thing** — a string where a number was expected, a list where a dataframe was expected
- Inspect first: `type(x)`, `df.dtypes`

Print, log, inspect — before guessing

- Guessing is editing code you have not looked at, to fix data you have not looked at
- One `print()` right above the failing line beats ten minutes of theorising
- Print the thing the line uses: `print(df.columns)`, `print(type(x))`,
`print(path.exists())`
- Then run it again, with the same command that failed
- Remove the print when done; a diagnosis note is where the finding lives

Minimal reproducible example

```
import pandas as pd
df = pd.read_csv("data/raw/sales.csv")
print(df.columns.tolist())
print(df["revenue"].sum())          # the line that fails, isolated
```

- Strip everything that is not needed to produce the error
- If four lines still fail, the bug is in those four lines
- If four lines work, the bug is in what you removed — bisect
- This is also the best thing to paste when asking for help

How to ask for help well

A question a TA, a colleague, or an assistant can actually answer contains:

1. The **exact** error, pasted, not paraphrased
2. The **failing line** and the file it lives in
3. Where things are: file tree, working directory, the command you ran
4. What you **expected** and what **happened** instead
5. What you **already tried** — and what it showed you

"It doesn't work" contains none of these.

Git thread — what happened before?

git log — the history of your checkpoints

```
git log --oneline
```

```
e7c21a4 Rename revenue column after load  
b93f0d2 Add monthly_total()  
a41f9c2 Initial working report
```

- Every commit is a moment the project was in a known state
- Small commits mean small diffs between working and broken

If it worked before, history tells you what changed

```
git log --oneline -5 -- scripts/report.py  
git diff a41f9c2 -- scripts/report.py
```

- Which commits touched this file? That is the suspect list
- What is different between the last working commit and now? That is the diff
- You are not guessing what you changed; you are **reading** it
- Lab 5 ships with a small history — use it

AI thread — a prompt made of evidence

"Fix this" versus evidence

Bad prompt

```
my script doesn't work fix this
```

Good prompt

```
Running `uv run python scripts/report.py` from the project root gives  
KeyError: 'revenue' at scripts/report.py line 8: df["revenue"].sum().  
print(df.columns.tolist()) shows ['date', 'Revenue', 'region'].  
I expected a column named 'revenue'. I have not changed the CSV.  
What are the most likely causes, and what should I inspect next?
```

AI gets dramatically better when you provide evidence instead of frustration

- The five ingredients from "asking well" are the prompt
- Ask for **likely causes and what to inspect**, not "fix it"
- Try **one** change; run the **same** command; read the result yourself
- Then explain the fix in your own words — that sentence goes in the diagnosis note
- Your AI-use note names what you asked and how you verified the answer

5-minute stand-and-stretch

Stand up. Leave the laptop. Back in five for Lab 5.

Lab 5 — The traceback is the map

First 15 minutes: no AI. The traceback and the failing file only.

- Chat windows, AI editors, assistant tabs: **closed**. Search engines are fine
- Why: the reflex we are installing is "*check the columns*", not "*ask what to check*"
- Only one of those reflexes survives a closed-book exam
- Solve it early? Go straight to checkoff. The window is a floor on struggle, not a ceiling on speed
- After 15 minutes, normal AI policy resumes

The lab: four failures in sequence

`scripts/report.py` fails **four times**. Each fix reveals the next.

1. A wrong **filename**
2. A wrong **column name**
3. A value that **will not convert**
4. A function called with the **wrong kind of input**

You already know which error type each one raises. Prove it.

Run it, diagnose it, note it

```
uv sync  
uv run python scripts/report.py
```

- Run it **exactly like that**, from the project root, every time
- For each failure: error type, failing line, what you inspected, what you changed, how you verified
- **One diagnosis note per failure** — four five-part notes in `DIAGNOSIS.md` ; part 3 quotes the two anchors
- Commit after each fix; `git log` should show four checkpoints
- **Checkoff**: a TA picks one failure; you explain it in 60–90 seconds, no notes